

Elastic-BLAST exercises

For these exercises, you search sequences from the chicken gut metagenome against a well curated 16S database for bacteria and archaea. You produce the BLAST results in a tabular format that allows you to calculate some specific values for each query-subject match. From this information, you can infer some properties of the query sequences such as a putative taxonomic classification or whether the query is a partial sequence.

The 16S database consists of 16S ribosomal RNA sequences that correspond to bacteria and archaea type materials. This database is part of the RefSeq Targeted Loci Project. See <https://www.ncbi.nlm.nih.gov/refseq/targetedloci/> for more information.

The chicken gut metagenome is from the WGS project KCET01 (<https://www.ncbi.nlm.nih.gov/Traces/wgs/KCET01?display=contigs&page=1>). This project contains nearly 4 million sequences, but here we look at only the first 299,000 sequences.

Definition of a “good match”:

A query has a good match to the 16S_ribosomal_RNA database if these statements are true:

- 1.) The match consists of one alignment
- 2.) The match shows an 80% identity between the query and database sequence.
- 3.) In the match, the database sequence covers (aligns to) at least 90% of the query. The qcovhsp value in the tabular output specifies the coverage.

Tasks:

1.) Login to a GCP instance and setup Elastic-BLAST:

- You will need to install kubectl (client for k8s). See <https://elbdoc.readthedocs.io/en/latest/requirements.html> for the kubectl link.
- You will need to install the elastic-blast client. See <https://elbdoc.readthedocs.io/en/latest/quickstart-gcp.html> for instructions.
- You may need to update python3-distutils. If so, use:
 - `sudo apt-get install -y python3-distutils`
- If you are creating a new instance, Ubuntu 18.04 with at least a 100G of disk space is suggested.

2.) Run Elastic-BLAST using the config file blastn.16S.ini. Notes:

- You will need to edit the file blastn.16S.ini to add your results bucket. Use the bucket you created earlier and add “/results” to the end.
- Elastic-BLAST will construct a default cluster name for you (elasticblast-`{USER}`). Other values in the cluster section of the configuration file have been optimized for the 16S database you will be searching.
- The query file is in a cloud bucket and can be accessed directly by Elastic-BLAST.
- The tabular format is specified as '7 std qcovhsp staxid ssciname'. Slide 15 of the presentation lists the fields that “std” specifies.
- The Elastic-BLAST search should take about 37 minutes.

3.) Identify the good matches in your results and produce a file listing, for each match, query ID, subject ID, genus. From this:

- Find the genus that occurs most often per query sequence, or
- Find the genera of the top five matches per query sequence, or
- Use some other method of your choosing to identify a genus for the query.

4.) Produce a list of the query sequences **without** good matches.

5.) Search a few of the query sequences from 4.) against nt using the BLAST webpage to see if you can identify them.

